

Memory management, File IO, ...

Jamal Hussein

jah1g12@ecs.soton.ac.uk

Electronics and Computer Science

December 4, 2015

Outline

1 Memory layout - Heap


2 Memory Allocation

3 File IO

4 Miscellaneous

Memory Layout - Heap

- Global variables, static variables and program instructions get their memory in permanent storage area
- local variables are stored in Stack
- The memory space between these two region is known as Heap area.
- This region is used for dynamic memory allocation during execution of the program.
- The size of heap keep changing.



Memory Allocation (stdlib.h)

- The header `<stdlib.h>` declares functions for storage allocation
- `malloc` returns a pointer to space for an object of a specified size, or `NULL` if the request cannot be satisfied. The allocated space is uninitialized

```
void *malloc(size_t size)
```

- `free` deallocates the space that previously allocated by `calloc`, `malloc`, or `realloc`

```
void free(void *P)
```

- `size_t` is an unsigned integer defined in `stddef.h` as (on my computer):

```
typedef long unsigned int size_t;
```

Memory Allocation (malloc)

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main() {
    char *language;

    language = malloc( 200 * sizeof(char) );

    if( language == NULL ) {
        fprintf(stderr, "Error - ...\n");
    }
    else {
        strcpy( language, "C programming language");
    }
    printf("Language: %s\n", language );
    free(language);
    return 0;
}
```

Memory Allocation (stdlib.h)

- `calloc` returns a pointer to space for an array of objects, each of size `size`, or `NULL` if the request cannot be satisfied. The space is initialized to zero bytes.

```
void *calloc(size_t nobj, size_t size)
```

- `realloc` changes the size of the previously allocated memory by a new size.
- `realloc` returns a pointer to the new space, or `NULL` if the request cannot be satisfied, in which case the old allocated memory is unchanged.

```
void *realloc(void *p, size_t size)
```

Memory Allocation (calloc)

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>

int main() {
    char *language;

    language = calloc( 200, sizeof(char) );

    if( language == NULL ) {
        fprintf(stderr,
            "Error - ...\n");
    }
    else {
        strcpy( language, "C programming language");
    }
    printf("Language: %s\n", language );
    free(language);
    return 0;
}
```

Memory Allocation (realloc)

```
int main() {  
    char *language;  
    language = malloc( 25 * sizeof(char) );  
    if( language == NULL ) {  
        fprintf(stderr, "...\\n");  
    }  
    else {  
        strcpy( language, "C programming language");  
    }  
    language = realloc( language, 10 * sizeof(char) );  
    if( language == NULL ) {  
        fprintf(stderr, "...\\n");  
    }  
    else {  
        strcat( language, " tutorial");  
    }  
    printf("Language: %s\\n", language );  
    free(language);  
    return 0;  
}
```


Memory Allocation (pointer)

```
int main() {  
    int *iptr = malloc (sizeof(int));  
    float *fptr = malloc (sizeof(float));  
    double *dptr = malloc (sizeof(double));  
  
    *iptr = 100;  
    *fptr = 4.13f;  
    *dptr = 459000.0;  
  
    printf(" *iptr: %d\n", *iptr);  
    printf(" *fptr: %.2f\n", *fptr);  
    printf(" *dptr: %.2f\n", *dptr);  
  
    free(iptr);  
    free(fptr);  
    free(dptr);  
    return 0;  
}
```

File IO (stdlib.h)

- `fopen` opens the named file, and returns a stream, or `NULL` if the attempt fails

```
FILE *fopen(const char *filename, const  
char *mode)
```

Legal values for mode include:

"r"	open text file for reading
"w"	create text file for writing; discard previous contents if any
"a"	append; open or create text file for writing at end of file
"r+"	open text file for update (i.e., reading and writing)
"w+"	create text file for update; discard previous contents if any
"a+"	append; open or create text file for update, writing at end

File IO (stdio.h)

- `fflush` causes any buffered but unwritten data to be written on an output stream

```
int fflush(FILE *Stream)
```

- on an input stream, the effect is undefined
 - It returns `EOF` for a write error, and zero otherwise
 - `fflush(NULL)` flushes all output streams
 - called between a write and a read
 - `fclose` flushes any unwritten data for stream, discards any unread buffered input
- ```
int fclose(FILE *Stream)
```
- frees any automatically allocated buffer, then closes the stream
  - It returns `EOF` if any errors occurred, and zero otherwise

# File IO

## ■ Writing a file:

- `fputc()` writes the character value to an output stream

```
int fputc(int c, FILE *fp);
```

- `fputs()` writes a string to an output stream

```
int fputs(const char *s, FILE *fp);
```

- `fprintf` write a string to an output stream

```
int fprintf(FILE *fp, const char *format,
...)
```

# File IO

```
#include <stdio.h>

int main() {
 FILE *fp;

 fp = fopen("test.txt", "w+");

 fprintf(fp, "This is printed by fprintf...\n");
 fputs("This is printed by fputs...\n", fp);

 fclose(fp);
}
```

# File IO

## ■ Reading a file:

- `fgetc()` reads a character from the input file. The return value is the character read, or in case of any error, it returns `EOF`

```
int fgetc(FILE * fp);
```

- `fgets()` reads up to `n-1` characters from an input stream. It copies the read string into a buffer, appending a null character to terminate the string.

```
char *fgets(char *buf, int n, FILE *fp);
```

- `fscanf` reads strings from a file, but it stops reading after encountering the first space character

```
int fscanf(FILE *fp, const char *format,
...)
```

# File IO

```
#include <stdio.h>
main() {
 FILE *fp;
 char buff[255];
 fp = fopen("test.txt", "r");

 fscanf(fp, "%s", buff);
 printf("1: %s\n", buff);

 fgets(buff, 255, (FILE*)fp);
 printf("2: %s\n", buff);

 fgets(buff, 255, (FILE*)fp);
 printf("3: %s\n", buff);

 fclose(fp);
}
```

# File IO

This is the output of the program:

```
1: This
2: is printed by fprintf...
3: This is printed by fputs...
```



# Command-line Arguments

- In C it is possible to accept command line arguments
- Command-line arguments are given after the name of a program
- passed to the program by the operating system
- The full declaration of main looks like this:  

```
int main (int argc, char *argv[])
```

  - `argc`: number of command line arguments including the name of the program
  - `argv`: list of all command line arguments
    - `argv[0]`: the program name (or an empty string)
    - `argv[1]` to `argv[argc-1]`: the actual command line arguments

# Command-line Arguments

```
#include <stdio.h>
int main (int argc, char *argv[]) {
 if (argc != 2) {
 printf("usage: %s filename", argv[0]);
 }
 else {
 FILE *file = fopen(argv[1], "r");
 if (file == NULL) {
 printf("Could not open file\n");
 }
 else {
 int x;
 while((x = fgetc(file)) != EOF) {
 printf("%c", x);
 }
 fclose(file);
 }
 }
}
```

# Variable Arguments

- Sometimes, you would like to have a function that accept an arbitrary number of arguments
- for example a function that accepts any number of values and returns the average  

```
double average(int num, ...)
```
- the last argument is written as ellipses, i.e. three dots (...)
- the one just before the ellipses is always an `int` which will represent the total number variable arguments passed

# Variable Arguments

- The header file `stdarg.h` provides the functions and macros to implement the functionality of variable arguments as follows
  - Define a function with its last parameter as ellipses and the one just before the ellipses is always an `int` which will represent the number of arguments.
  - Create a `va_list` type variable in the function definition
  - Use `int` parameter and `va_start` macro to initialize the `va_list` variable to an argument list
  - Use `va_arg` macro and `va_list` variable to access each item in argument list
  - Use a macro `va_end` to clean up the memory assigned to `va_list` variable

# Variable Arguments

```
#include <stdio.h>
#include <stdarg.h>
double average(int num,...) {
 va_list valist;
 double sum = 0.0;
 int i;
 va_start(valist, num);
 for (i = 0; i < num; i++) {
 sum += va_arg(valist, int);
 }
 va_end(valist);
 return sum/num;
}
int main() {
 printf("Average of 2, 3, 4, 5 = %f\n",
 average(4, 2,3,4,5));
 printf("Average of 5, 10, 15 = %f\n",
 average(3, 5,10,15));
}
```

# typedef

- The `typedef` keyword allows the programmer to create new names for types such as `int`
- Typedefs can be used both to provide more clarity to your code and to make it easier to make changes to the underlying data types

```
typedef long unsigned int size_t;
```

- In C, `struct` variables must be declared by a combination of the keyword `struct` and the name of the struct

```
typedef struct Books
```

```
...
```

```
} Book;
```

# typedef

```
#include <stdio.h>
#include <string.h>
typedef struct employee {
 char name[50];
 int salary;
} emp ;
void main() {
 emp e1;
 printf("Enter Employee record\n");
 printf("Employee name:\t");
 scanf("%s", e1.name);
 printf("Enter Employee salary:\t");
 scanf("%d", &e1.salary);
 printf("\nstudent name is %s\n", e1.name);
 printf("roll is %d\n", e1.salary);
}
```

# typedef vs #define

- `#define` is a C-directive which is also used to define the aliases for various data types similar to `typedef` but with the following differences:
    - `typedef` is limited to giving symbolic names to types only whereas `#define` can be used to define alias for values as well
- ```
#define TRUE 1
#define FALSE 0
```
- `typedef` interpretation is performed by the compiler whereas `#define` statements are processed by the pre-processor

Preprocessors

- Preprocessors are a way of making text processing with your C program before they are actually compiled
- Before the actual compilation of every C program it is passed through a Preprocessor.
- The Preprocessor looks through the program trying to find out specific instructions called Preprocessor directives that it can understand.
- All Preprocessor directives begin with the # (hash) symbol

Preprocessors

<code>#define</code>	Substitutes a preprocessor macro.
<code>#include</code>	Inserts a particular header from another file.
<code>#undef</code>	Undefines a preprocessor macro.
<code>#ifdef</code>	Returns true if this macro is defined.
<code>#ifndef</code>	Returns true if this macro is not defined.
<code>#if</code>	Tests if a compile time condition is true.
<code>#else</code>	The alternative for <code>#if</code> .
<code>#elif</code>	<code>#else</code> and <code>#if</code> in one statement.
<code>#endif</code>	Ends preprocessor conditional.
<code>#error</code>	Prints error message on stderr.
<code>#pragma</code>	Issues special commands to the compiler, using a standardized method.

Preprocessors

- `#define` is used to define values or macros that are used by the preprocessor to manipulate the program source code
- By convention, values defined using `#define` are named in uppercase
- can produce quite unexpected results if not done right

```
#define MAX(a,b) ((a)>(b)?(a):(b))
```

```
:
```

```
i = 2;
```

```
j = 3;
```

```
k = MAX(i++, j++);
```

- the definition must parenthesize every use of `a,b` in the macro definition

Preprocessors

```
#define SLICES 8
#define ADD(x) ( (x) / SLICES )
int main(void) {
    int a = 0, b = 10, c = 6;

    a = ADD(b + c);
    printf("%d\n", a);
    return 0;
}
```

- The `#undef` directive undefines a macro. The identifier need not have been previously defined.
- `#ifdef` and `#ifndef` are used to check whether the code block following them are defined or undefined respectively

Preprocessors

- ANSI C defines some useful preprocessor macros and variables

<code>__FILE__</code>	The name of the current file, as a string literal
<code>__LINE__</code>	Current line of the source file, as a numeric literal
<code>__DATE__</code>	Current system date, as a string
<code>__TIME__</code>	Current system time, as a string
<code>__STDC__</code>	defined as 1 when the compiler complies with the ANSI standard.

Preprocessors

```
#include <stdio.h>
main() {
    printf("File  :%s\n", __FILE__ );
    printf("Date  :%s\n", __DATE__ );
    printf("Time  :%s\n", __TIME__ );
    printf("Line  :%d\n", __LINE__ );
    printf("ANSI  :%d\n", __STDC__ );
}
```

The output:

```
File  :macros2.c
Date  :Dec  3 2015
Time  :16:41:02
Line  :8
ANSI  :1
```

Header files

- A header file is a file with extension `.h` which contains C function declarations and macro definitions to be shared between several source files.
- There are two types of header files:
 - the files that the programmer writes
`#include <stdio.h>`
 - the files that comes with your compiler
`#include "my_header.h"`

Header files

```
/* my_header.h file */  
int max (int a, int b);
```

```
/* my_header.c file */  
int max( int x, int y){  
    return x > y ? x : y;  
}
```

```
/* main program */  
#include <stdio.h>  
#include "my_header.h"  
int main () {  
    printf("max(2,3) = %d\n", max (2, 3));  
    return 0;  
}
```


Header files (Once-Only Headers)

- If a header file happens to be included twice, the compiler will process its contents twice and it will result in an error
- The standard way to prevent this is to enclose the entire real contents of the file in a conditional pre-processing statement:

```
#ifndef MY_HEADER
#define MY_HEADER

int max (int a, int b);

#endif
```